

Mohsen Mostafa

AI Team Lead – AI / ML & Data Infrastructure

Cairo, Egypt | +20 102 556 5499 | mohsen.mostafa.ai@outlook.com | [Portfolio](#)

PROFESSIONAL SUMMARY

- AI Team Lead & Research Engineer with 8+ years in AI/ML/CV/GenAI and 10+ years in IT infrastructure.
- Proven track record of production deployments with measurable business impact.
- Built end-to-end MLOps pipelines (Kafka, Spark, Hadoop, MLflow, Kubeflow) for real-time Robotics & IoT anomaly detection.

CORE SKILLS

Languages & Data	Python, C++, Pandas, SQL, PySpark
ML / DL Frameworks	PyTorch, TensorFlow, XGBoost, scikit-learn, Isolation Forest, LSTM Autoencoders
GenAI / LLM	Hugging Face, Falcon, DeepSeek, Llama, LangChain, RAG, ColBERT
Computer Vision	YOLOv8, OpenCV, Mask2Former
Streaming & Big Data	Apache Kafka, Spark Structured Streaming, Hadoop HDFS, Parquet
MLOps & DevOps	MLflow, Kubeflow, Airflow, Docker, Kubernetes, KServe, FastAPI, GitLab CI
Databases	PostgreSQL, MongoDB, Redis, Pinecone
Cloud & Infra	AWS SageMaker/EC2, GCP Vertex AI, Prometheus, Grafana, ONNX, TensorRT

PROFESSIONAL EXPERIENCE

AI Team Lead | [GraviLog \(Remote\)](#) Sep 2025 – Present

- End-to-end AI system architecture for multimodal CV+LLM (TensorRT/ONNX) → 40% latency reduction.
- Defined technical roadmap for 3 products; delivered C-level presentations.
- Led RAG retrieval (Pinecone + embedding); mentored 4 engineers in MLOps & system design.

Senior AI Developer | [OMC Technologies \(Remote\)](#) May 2023 – Aug 2025

- Built AI agents (Falcon, DeepSeek, Qwen, Llama) → automated 3 workflows, 40% manual effort reduction.
- Production RAG (LangChain + Pinecone + Falcon/DeepSeek) → 35% hallucination reduction.
- Multimodal healthcare AI → F1 ↑19%, latency 210ms → 145ms (30% via streaming + FP16).
- Fraud detection AUC ↑17%, false positives ↓22% (FastAPI, p99 ≤ 95ms @ 500 RPS).
- Managed AI infrastructure with 24 GPUs (12 nodes) + 15 CPU nodes for real-time IoT anomaly detection; optimized distributed training and inference using TensorRT, Kubeflow, and KServe.
- Mentored 5 juniors → 2 promotions.

AI Developer & Researcher | [Nexgro \(South Africa\)](#) Apr 2020 – Apr 2023

- Grain quality CV pipeline → 98% accuracy across 6 types.
- Drone/satellite anomaly detection (NDVI + YOLOv8) → manual inspection efficiency ↑160%.
- LLM RAG assistant on technical manuals; document automation saved 15 person-hours/week.
- RL logistics prototype → travel distance ↓12%.

Data Scientist | [Nexgro](#) Mar 2018 – Mar 2020

- Forecast models → logistics costs ↓18%.
- Predictive maintenance prevented 3 breakdowns, saving ~\$45K.
- Inventory optimization → spoilage ↓18%.

Systems Administrator | [Egyptian Co.](#) Jun 2013 – Mar 2018

- IT infrastructure management for 200+ users → 99.8% uptime; multi-site network & servers.
- Automation scripts reduced manual ticket handling by 25%.
- Self-directed pivot to data science (certifications 2016–2018); managed LAN/WAN, firewalls, backup for 150+ nodes.

IT Network Administrator | [International Co. for Trade](#) Aug 2009 – Jun 2013

- Administered LAN/WAN, firewalls, and backup solutions for 150+ nodes.

IT Network Specialist | [International Co. for Trade](#) Aug 2005 – Jun 2009

- First-line network support, hardware procurement, and vendor coordination.

EDUCATION

B.Sc. Information Systems
High Inst. for Specialized Technological Studies, Egypt
2000 – 2005

LANGUAGES

Arabic (Native) | English (Fluent – C1/C2)

RESEARCH PREPRINTS

- Geometric Deep Learning for Point Cloud Segmentation – under revision for WACV 2027.
- Physics-Informed Neural Networks for Fluid Flow – arXiv 2025.

Full list on *ORCID*: 0009-0004-4478-0317

CERTIFICATIONS

- Deep Learning Specialization – DeepLearning.AI
- TensorFlow Advanced Techniques – Google
- Google Cloud MLOps
- Total Data Quality – University of Michigan

LEADERSHIP & MENTORING

- Mentored 5+ engineers (2 promotions); led cross-functional AI initiatives.
- Conducted internal LLM/MLOps workshops.
- Designed full MLOps showcase: Kafka → Spark Streaming → HDFS → MongoDB → PostgreSQL → MLflow → Kubeflow → FastAPI for real-time predictive maintenance.